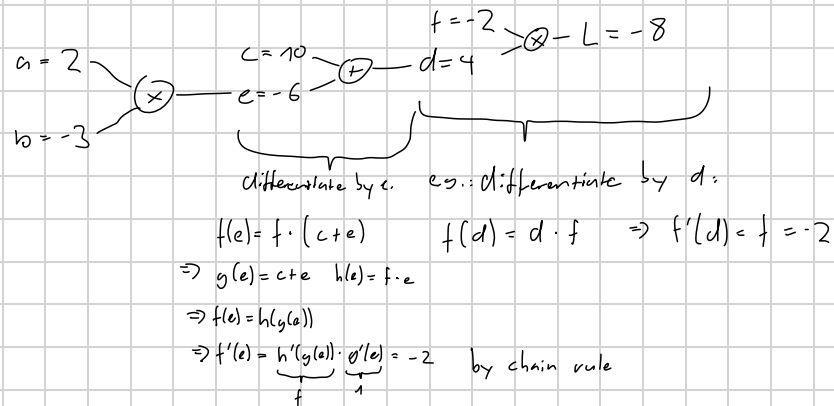


# Backpropagation


↳ figure out how much changing certain inputs changes outputs

Example:




Like this we can backpropagate through the entire network and figure out the gradients by multiplying the local gradients by global previous (from backward direction) gradients.

For example above:




## Single neuron



forward pass: calculate output depending on inputs  $\rightarrow$  prediction  
backward pass: backpropagate and find gradients  $\rightarrow$  learning

MLP: Multi-layer-perception



each layer fully connected to previous layer

```
class Neuron:
    def __init__(self, nIn):
        self.w = [Value(random.uniform(-1, 1)) for _ in range(nIn)]
        self.b = Value(random.uniform(-1, 1))
    def __call__(self, x):
        # calculate w * x + b
        act = sum(xi * wi for xi, wi in zip(self.w, x))
        out = act.tanh()
        return out

class Layer:
    def __init__(self, nIn, nOut):
        self.neurons = [Neuron(nIn) for _ in range(nOut)]
    def __call__(self, x):
        outs = [n(x) for n in self.neurons]
        return outs

class MLP:
    def __init__(self, nIn, nOut):
        sz = [nIn] + nOut # list with n inputs, nOut is list with outputs per level
        self.layers = [Layer(sz[i], sz[i+1]) for i in range(len(nOut))]
    def __call__(self, x):
        for layer in self.layers:
            x = layer(x)
        return x

x = [2.0, 3.0, -1.0]
n = MLP(3, [4, 4, 1])
n(x)
```

3 inputs, hidden layers: 4, 4, output: 1

## Gradient descent

through backpropagation we figured out the gradients of each parameter (=weights + biases of neurons in MLP). In gradient descent we then set data to  $\text{data} = \text{step-size} \cdot \text{grad}$ , we then iteratively recalculate gradients and adjust the parameters slightly. Through this we try to find the minimum of parameters for loss.


```
#gradient descent
# reset all gradients to zero
for k in range(20):

    # forward pass
    ypred = [n(x) for x in xs]
    loss = sum((yi - predi)**2 for yi, predi in zip(ys, ypred))

    # backward pass
    for p in n.parameters():
        p.grad = 0.0 #need to reset since otherwise backwards accumulates => huge step size
    loss.backward()

    # update
    for p in n.parameters():
        p.data += -0.05 * p.grad #step size 0.05
```

Visually:



Batch: On tiny networks we can go through all data in the forward pass, however for actual useful networks, the training data is too large  $\Rightarrow$  we take random subset of the data and use that for forward pass.

Learning rate decay: As the network stabilizes, we decrease learning rate per step to avoid overshooting.